

Learning Objectives

- Describe the use of Unix I/O library functions.
- Understand when short reads and writes happen and how to handle them.
- Recognize the implementation differences and resulting performance differences between Unix I/O and Standard I/O.
- Predict the behavior of multi-process programs when file descriptors are opened and copied (with `dup` and `dup2`).
- Describe how shell I/O redirection is implemented in terms of Unix I/O operations.

Getting Started

The directions for today's activity are on this sheet, but they refer to programs that you'll need to download. To get set up, run these commands on a shark machine:

```
$ wget http://www.cs.cmu.edu/~213/activities/system-io.tar
$ tar xf system-io.tar
$ cd system-io
$ make
```

1 Unix I/O

The Unix I/O API consists of low-level functions for opening, closing, reading from, and writing to files. Open files are represented in this API by *file descriptors*, which are non-negative integers.

The file descriptors 0, 1, and 2 are already open when a program starts up. They correspond to the Standard I/O streams standard input (`stdin`), standard output (`stdout`), and standard error (`stderr`) respectively.

Examine the program `unixcat.c`, which opens a file, reads its contents, and writes its contents to `stdout`.

Problem 1. When you're ready, try using this program to print its own source:

```
$ ./unixcat unixcat.c
```

What went wrong? Edit `unixcat.c` to fix the problem (hint: it's on the line after the `FIXME` comment); run `make` again to recompile, then try using it again. What did you have to do to make it work correctly?

When the remaining bytes are smaller than the buffer size, the original code stops reading. It should still read until the EOF.

Using `read() > 0` condition resolves the issue.

2 Standard I/O

The C standard does not specify the Unix I/O functions. It does specify the *standard I/O* functions, which operate on files at a higher level. (These are implemented using the Unix I/O functions on Unix systems.) Some of the standard I/O functions are very similar to Unix I/O functions, such as `fread` and `read`. Others, such as `printf`, provide facilities for complex “formatted” I/O. All of the standard I/O functions represent open files with `FILE` objects, which are structs containing file descriptors but also various other data.

One of the most important reasons to use standard I/O is *buffering*: data can be temporarily stored inside each `FILE` (in a char array) to reduce the number of calls to read and write that are necessary. To see how buffering affects a program's output, look at the program `three.c`. It prints three three-word phrases using different combinations of `printf` and `write`. When you're ready, run the program with `./three`.

Problem 2. Which of the three phrases were printed in the same order as they appear in the source code? Why did this occur? What is different about the way the three phrases were printed that caused this?

The system call `write` will take effect straight away, while the standard `printf` will store the content in an internal buffer first, then flush it when specified by the `fflush` function.

In this example, the `printf` buffer is only cleared when hitting a new line.

Problem 3. You can use the `strace` program to examine the system calls made by `three`. Run it like this:

```
$ strace -e trace=write ./three > /dev/null
```

What do you notice about the calls to `write`? Does this agree with what you observed?

Yes. The calls to write are exactly the order we expected considering the buffer used in `printf`

Note: In this problem, `strace` was directed to report *only* calls to `write`, and I/O redirection was used to suppress the output directly from `three`. You might be curious to see what happens if you simply run

```
$ strace ./three
```

(This will produce a lot of output—be prepared to scroll back.)

2.1 Buffering and Performance

Buffered I/O aims to increase efficiency by reducing the number of calls to `read` and `write`, which have a lot of overhead (tens of thousands of clock cycles). Examine the program `timing.c`, which measures the time it takes to write data to `/dev/null` one byte at a time, using both Unix I/O (`write`) and Standard I/O (`fputc`). (`/dev/null` is a special file called the “null device,” or, colloquially, the “bit bucket,” which discards anything written to it. We’ve already used it once in this activity.)

Problem 4. The program `timing` takes one command-line argument, the number of bytes to write. When you are done looking over its code, try having it write just one byte each way:

```
$ ./timing 1
```

Which way is faster, `write` or `fputc`? Is that what you expected? You may need to run it several times to see a pattern.

The unix write is about twice faster than `fputc`.
Because the std library function `fputc` uses a buffer.
But for a single byte it is actually an overhead.

Problem 5. Try increasing the number of bytes written, in steps. How many bytes do you need to write before buffering is faster? If you make the number of bytes larger and larger, what do you think the *asymptotic* performance curves look like?

Buffering start to be worth it after 100 bytes
Both linear but write has a steeper gradient

Problem 6. Based on what you have just observed, when do you think you should

use Unix I/O functions, and when do you think it will be better to use Standard I/O?

Unix I/O function only works well when writing a small amount of data or large data all at once. However, standard I/O is better when handling a lot of reads and writes, e.g. process a file character by character or line by line

Problem 7. (advanced) Edit `timing.c`. On the line that reads

```
setvbuf(stream, NULL, _IOFBF, 0); // Buffer in large chunks.
```

change `_IOFBF` to `_IOLBF`. Recompile (`make`). Run `timing` again, with a large argument (say, 100000). What changed? Why might that have happened?

3 File Descriptors, Fork, and Dup2

Each time the `open` function is called, a new *open file table entry* is created and the file descriptor corresponding to that entry is returned. However, the `dup` and `dup2` functions can be used to *duplicate* a file descriptor. Duplicated file descriptors point to the *same* open file table entry. `dup` takes the old file descriptor as an argument and returns the duplicate. `dup2` takes both the old and new descriptors as arguments, and makes the new descriptor be a duplicate of the old. If the new descriptor was open already, it is closed first (atomically).

Examine the program `doublecat.c`, paying particular attention to the `print2` function, which takes in two file descriptors and prints each file, one character at a time.

Problem 8. The file `abcde.txt` contains the characters `abcde`, followed by a newline character. If you were to run `$./doublecat abcde.txt` what would be printed in each case?

case 1: a b c d e

case 2: a a b b c c d d e e

case 3: a b c d e

Problem 9. Run the program as suggested (`$./doublecat abcde.txt`). Were your predictions correct? Did the output differ in the three cases? Why do you get the output you do?

In case 1 and 3 there is only one `open()` operation -> one file position

To further complicate matters, child processes share the open file descriptors of their parents. (It is as if `fork` calls `dup` for each descriptor—but the new descriptors go into the child's file descriptor table and have the same numbers that they did in the parent.)

Examine the program `childcat.c`, which forks two child processes, each of which print two letters from a shared file descriptor, while the parent prints one letter from that file.

Problem 10. If you run `childcat` on `abcde.txt`, what could be printed? Take a moment and write your guess below. Then, run `$./childcat abcde.txt`. Did the output match what you expected?

abcde

Problem 11. Run `$./childcat abcde.txt` several more times. Does it *always* print the same thing? Look at the code carefully. Is it *guaranteed* to print the same thing always? If it isn't, are there any constraints on what it can and cannot print?

Not always same thing.

Child and grandchild will each print two characters in the order that appears in the file.

But the order of parent, child and grandchild are non-deterministic

4 Shell I/O Redirection

Shell I/O redirection is a tool for defining the input and output of shell commands in a high-level, uniform manner.

Problem 12. Try running the command

```
$ /bin/echo 15213 rocks > phrase.txt
```

This writes the string "15213 rocks" to the file `phrase.txt`, creating it if it does not exist. What Unix I/O and process control functions do you expect the shell will use to run this command?

fork(): for child -> open file("phrase.txt") -> dup(file as stdio) -> exec(echo)
for parent -> waitpid()

Problem 13. We can use `strace` to observe how the shell runs this command. Run it like this:

System I/O

```
$ strace -f -e trace=process,open,dup2,write \  
    /bin/sh -c '/bin/echo 15213 rocks > phrase.txt'
```

Does the series of operations you observe, match what you expect?

Most of it.

Problem 14. Examine the `phrase.txt` file. Did the traced invocation change its contents? Now try the following command (untraced):

```
$ /bin/echo 15213 rocks >> phrase.txt
```

What did that do to the file? What do you think the shell did differently because you used `>>`?

The shell appends the content to the file instead of overwriting it.

The flag in the `open` is `O_APPEND` instead of `O_TRUNC`.

In addition to redirecting standard input and output to/from files, the shell has the ability to *pipe* the output of one program to the input of another. You request this with the `|` operator.

Problem 15. An example of a command that uses pipes is

```
$ ps aux | grep $USER
```

This runs `ps aux`, which prints out a report on all the running processes on the current machine, and sends the output to the input of `grep $USER`, which searches for lines containing your username and prints only those lines. The output of `grep` goes to your terminal. (`$USER` is a *shell variable*. The shell will replace it with your actual username when it runs the `grep` program.)

To run this “pipeline”, what Unix I/O and process control functions will the shell call, and in what order? How does this differ from the calls done for simple redirection to file? (One of the functions that’s needed, we haven’t mentioned at all yet. `man 2`

System I/O

pipe.)

Fork two processes. One for ps and the other for grep.
The pipe function sets up the data channel between them.
It links the write fd from ps to read fd from grep. The shell then calls `execve` program.